Print index

Physical page numbers of Mechanistic_Interpretability_Complete_Anthology.pdf (page 1 = title page).
Print range = cover + original paper; type it into your print dialog. Text pages = original paper only.

NO.	PAPER	COVER	TEXT PAGES	PRINT RANGE	PAGES
A	**TRANSFORMER FOUNDATIONS AND CIRCUITS**				section 11–370
01	● Attention Is All You Need	13	14–28	**13–28**	16
02	● Zoom In: An Introduction to Circuits	29	30–59	**29–59**	31
03	○ Multimodal Neurons	61	62–110	**61–110**	50
04	● Mathematical Framework for Transformer Circuits	111	112–174	**111–174**	64
05	● In-context Learning and Induction Heads	175	176–236	**175–236**	62
06	○ Softmax Linear Units	237	238–279	**237–279**	43
07	● Toy Models of Superposition	281	282–343	**281–343**	63
08	● Interpretability in the Wild (IOI circuit)	345	346–370	**345–370**	26
B	**CAUSAL METHODS AND CIRCUIT DISCOVERY**				section 371–493
09	● Best Practices of Activation Patching	373	374–400	**373–400**	28
10	● Automated Circuit Discovery (ACDC)	401	402–436	**401–436**	36
11	○ Causal Abstractions of Neural Networks	437	438–463	**437–463**	27
12	○ Distributed Alignment Search (DAS)	465	466–493	**465–493**	29
C	**SUPERPOSITION AND SPARSE AUTOENCODERS**				section 495–904
13	○ Composition & Superposition	497	498–505	**497–505**	9
14	○ Language Models Can Explain Neurons	507	508–549	**507–549**	43
15	○ Finding Neurons in a Haystack	551	552–591	**551–591**	41
16	● SAEs Find Highly Interpretable Directions	593	594–613	**593–613**	21
17	● Towards Monosemanticity	615	616–719	**615–719**	105
18	● Scaling Monosemanticity	721	722–808	**721–808**	88
19	○ Scaling and Evaluating Sparse Autoencoders	809	810–843	**809–843**	35
20	○ JumpReLU Sparse Autoencoders *optional*	845	846–871	**845–871**	27
21	● A is for Absorption	873	874–904	**873–904**	32
D	**FEATURE CIRCUITS AND LARGE-SCALE INTERPRETABILITY**				section 905–1272
22	● Transcoders Find Interpretable Feature Circuits	907	908–936	**907–936**	30
23	● Circuit Tracing	937	938–1054	**937–1054**	118
24	● On the Biology of a Large Language Model	1055	1056–1188	**1055–1188**	134
25	● SAEBench	1189	1190–1231	**1189–1231**	43
26	● MIB: Mechanistic Interpretability Benchmark	1233	1234–1272	**1233–1272**	40
E	**THE 2026 FRONTIER**				section 1273–1392
27	◇ Prioritize Feature Consistency in SAEs	1275	1276–1315	**1275–1315**	41
28	● Circuits Are Sparse in the Neuron Basis	1317	1318–1359	**1317–1359**	43
29	◇ Weight-sparse Transformers	1361	1362–1392	**1361–1392**	32
F	**ALIGNMENT, STEERING, AND REASONING**				section 1393–1772
30	○ Representation Engineering	1395	1396–1451	**1395–1451**	57
31	○ Activation Addition (ActAdd)	1453	1454–1485	**1453–1485**	33
32	○ Refusal Is Mediated by a Single Direction	1487	1488–1527	**1487–1527**	41
33	● Unfaithful Chain-of-Thought Explanations	1529	1530–1561	**1529–1561**	33

Print index

Physical page numbers of Mechanistic_Interpretability_Complete_Anthology.pdf (page 1 = title page).
Print range = cover + original paper; type it into your print dialog. Text pages = original paper only.

NO.	PAPER	COVER	TEXT PAGES	PRINT RANGE	PAGES
34	○ Monitoring Monitorability	1563	1564–1653	**1563–1653**	91
35	◇ How Does Chain of Thought Think? *optional*	1655	1656–1670	**1655–1670**	16
36	○ Locate, Steer, and Improve (survey) *optional*	1671	1672–1772	**1671–1772**	102
G	**CONCLUSION**			section 1773–1857	
37	● Neural Networks Are Already Clean	1775	1776–1857	**1775–1857**	83

Bibliography and sources: pp. 1859–1865. Notes pages between papers are not part of any paper's range.

Section print ranges

A section range runs from its divider to the last page of its final paper (covers and notes pages included).
The second string prints only the papers (covers + text), skipping the divider and notes pages.

A Transformer foundations and circuits — 11–370

Papers 01–08 · 360 pages · papers only (no divider or notes pages):

13-59,61-279,281-343,345-370

B Causal methods and circuit discovery — 371–493

Papers 09–12 · 123 pages · papers only (no divider or notes pages):

373-463,465-493

C Superposition and sparse autoencoders — 495–904

Papers 13–21 · 410 pages · papers only (no divider or notes pages):

497-505,507-549,551-591,593-613,615-719,721-843,845-871,873-904

D Feature circuits and large-scale interpretability — 905–1272

Papers 22–26 · 368 pages · papers only (no divider or notes pages):

907-1231,1233-1272

E The 2026 frontier — 1273–1392

Papers 27–29 · 120 pages · papers only (no divider or notes pages):

1275-1315,1317-1359,1361-1392

F Alignment, steering, and reasoning — 1393–1772

Papers 30–36 · 380 pages · papers only (no divider or notes pages):

1395-1451,1453-1485,1487-1527,1529-1561,1563-1653,1655-1772

G Conclusion — 1773–1857

Paper 37 · 85 pages · papers only (no divider or notes pages):

1775-1857

Essential papers (P0)

20 papers · 1096 pages including covers

PRINT-DIALOG STRING FOR ALL P0 PAPERS

13-59,111-236,281-343,345-370,373-436,593-613,615-719,721-808,873-904,907-1231,1233-1272,1317-1359,1529-1561,1775-1857

NO.	PAPER	PRINT RANGE	PRINTED	READ	ANNOTATED
01	Attention Is All You Need	**13–28**	☐	☐	☐
02	Zoom In: An Introduction to Circuits	**29–59**	☐	☐	☐
04	Mathematical Framework for Transformer Circuits	**111–174**	☐	☐	☐
05	In-context Learning and Induction Heads	**175–236**	☐	☐	☐
07	Toy Models of Superposition	**281–343**	☐	☐	☐
08	Interpretability in the Wild (IOI circuit)	**345–370**	☐	☐	☐
09	Best Practices of Activation Patching	**373–400**	☐	☐	☐
10	Automated Circuit Discovery (ACDC)	**401–436**	☐	☐	☐
16	SAEs Find Highly Interpretable Directions	**593–613**	☐	☐	☐
17	Towards Monosemanticity	**615–719**	☐	☐	☐
18	Scaling Monosemanticity	**721–808**	☐	☐	☐
21	A is for Absorption	**873–904**	☐	☐	☐
22	Transcoders Find Interpretable Feature Circuits	**907–936**	☐	☐	☐
23	Circuit Tracing	**937–1054**	☐	☐	☐
24	On the Biology of a Large Language Model	**1055–1188**	☐	☐	☐
25	SAEBench	**1189–1231**	☐	☐	☐
26	MIB: Mechanistic Interpretability Benchmark	**1233–1272**	☐	☐	☐
28	Circuits Are Sparse in the Neuron Basis	**1317–1359**	☐	☐	☐
33	Unfaithful Chain-of-Thought Explanations	**1529–1561**	☐	☐	☐
37	Neural Networks Are Already Clean	**1775–1857**	☐	☐	☐